

Experiment - 3

Nikhil Kale
D15A/50

Aim - Apply Decision Tree and Random Forest for classification tasks

1. Dataset Source

Dataset used:

<https://www.kaggle.com/datasets/johnsmith88/heart-disease-dataset>

Title: Heart Disease Dataset

2. Dataset Description

The dataset contains medical attributes used to predict the presence of heart disease.

Number of records: 303

Number of features: 13 input features

Target variable: 1 output variable

Target Variable:

- 0 → No Heart Disease
- 1 → Heart Disease

Features include:

- age
- sex
- chest pain type (cp)
- resting blood pressure (trestbps)
- cholesterol (chol)
- fasting blood sugar (fbs)
- resting ECG (restecg)
- maximum heart rate (thalach)
- exercise induced angina (exang)
- oldpeak
- slope
- number of major vessels (ca)

- thal

Dataset Characteristics:

- All numeric features
- No missing values
- Binary classification problem

3. Mathematical Formulation

Decision Tree

Uses entropy or Gini index.

Gini Index:

$$Gini = 1 - \sum p_i^2$$

Entropy:

$$Entropy = - \sum p_i \log_2 p_i$$

Goal:

Select splits that minimize impurity.

Random Forest

Ensemble of Decision Trees.

Prediction:

$$\hat{y} = \text{Majority Voting of Trees}$$

Reduces variance and overfitting.

4. Algorithm Limitations

Decision Tree

- Prone to overfitting
- High variance
- Sensitive to noise

Random Forest

- Computationally expensive
- Less interpretable
- Large memory usage

5. Methodology / Workflow

1. Dataset loading
2. Data exploration
3. Train-test split
4. Feature scaling
5. Model training
6. Hyperparameter tuning
7. Performance evaluation
8. Model comparison

Workflow:

Dataset → Preprocessing → Split → Scaling → Train Model → Tune → Evaluate → Compare

6. Performance Analysis

Evaluation Metrics Used:

- Accuracy
- Precision
- Recall
- F1-score
- Confusion Matrix

Observations:

- Random Forest usually achieves higher accuracy
- Decision Tree may overfit

- Random Forest generalizes better

7. Hyperparameter Tuning

Decision Tree tuned parameters:

- max_depth
- min_samples_split
- criterion

Random Forest tuned parameters:

- n_estimators
- max_depth
- min_samples_split

GridSearchCV with 5-fold cross validation used.

Tuning improves generalization and prevents overfitting.

Code and Output -

```
[1] ✓ 2s
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

plt.style.use('seaborn-v0_8')
sns.set_palette("Set2")

[2] ✓ 0s
df = pd.read_csv('/content/heart.csv')
df.head()

  age sex cp trestbps chol fbs restecg thalach exang oldpeak slope ca thal target
  0 52 1 0 125 212 0 1 168 0 1.0 2 2 3 0
  1 53 1 0 140 203 1 0 155 1 3.1 0 0 3 0
  2 70 1 0 145 174 0 1 125 1 2.6 0 0 3 0
  3 61 1 0 148 203 0 1 161 0 0.0 2 1 3 0
  4 62 0 0 138 294 1 1 106 0 1.9 1 3 2 0
```

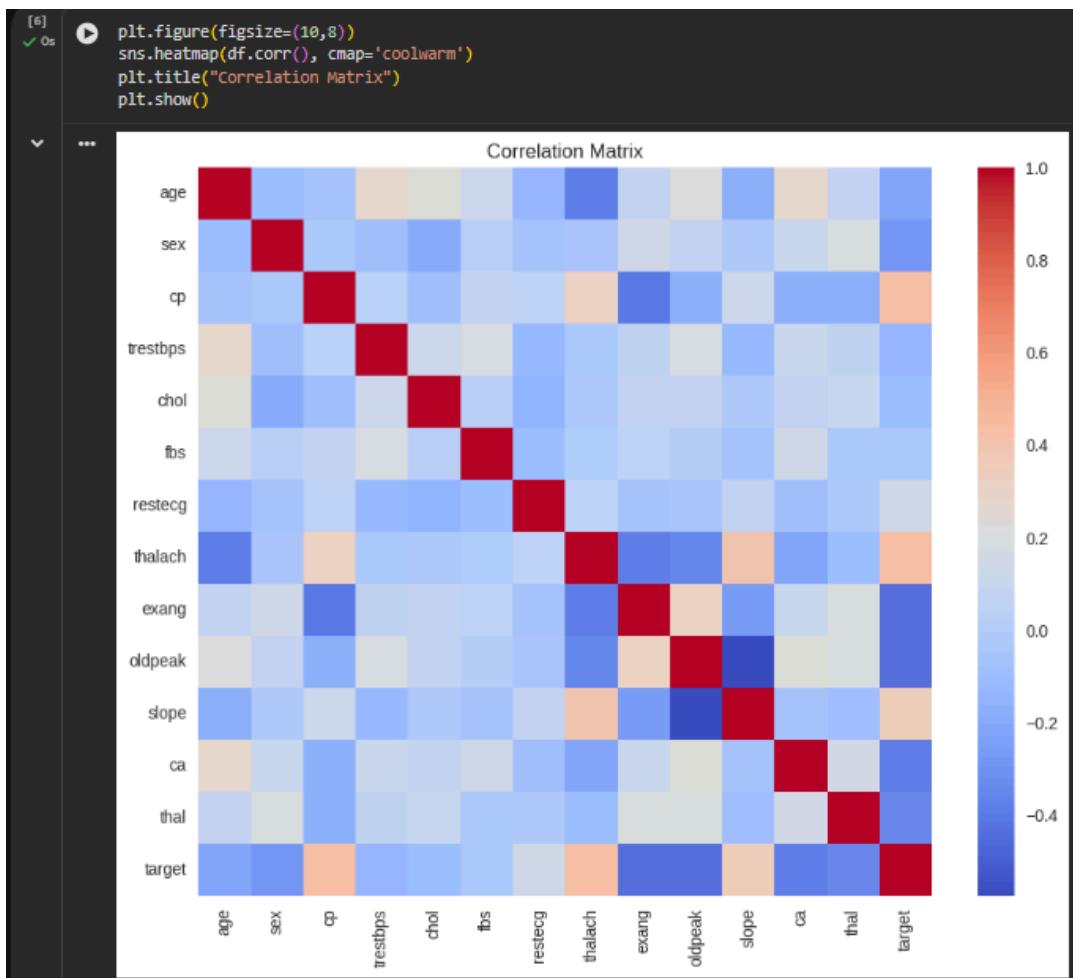
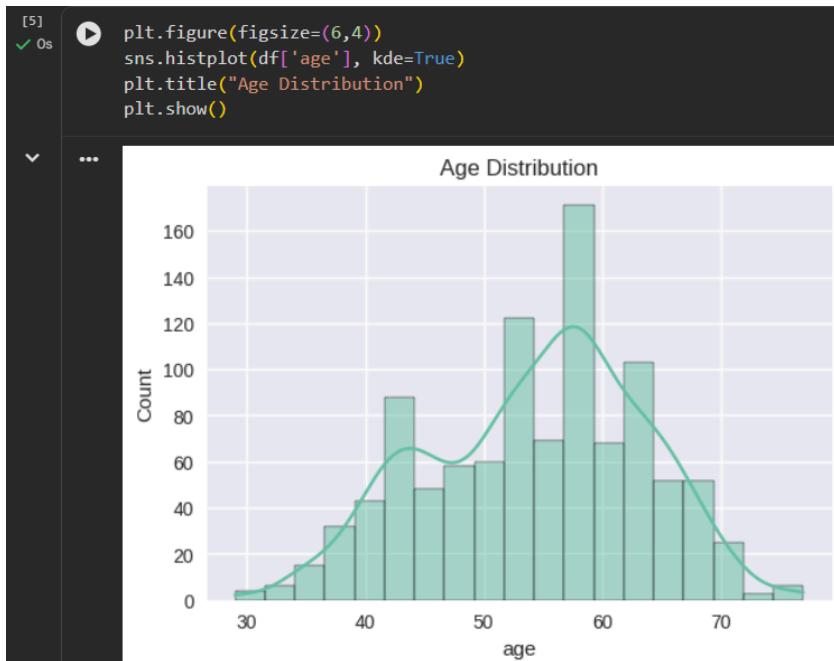
```
[3] df.info()
✓ 0s df.describe()
df.isnull().sum()

...
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   age        1025 non-null   int64  
 1   sex        1025 non-null   int64  
 2   cp         1025 non-null   int64  
 3   trestbps  1025 non-null   int64  
 4   chol       1025 non-null   int64  
 5   fbs        1025 non-null   int64  
 6   restecg   1025 non-null   int64  
 7   thalach   1025 non-null   int64  
 8   exang      1025 non-null   int64  
 9   oldpeak    1025 non-null   float64 
 10  slope      1025 non-null   int64  
 11  ca         1025 non-null   int64  
 12  thal       1025 non-null   int64  
 13  target     1025 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

0

	age	0
	sex	0
	cp	0







[8] ✓ 0s

```
X = df.drop('target', axis=1)
y = df['target']
```

[9] ✓ 0s

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

[10] ✓ 0s

```
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[11] ✓ 0s ⏪ dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)

y_pred_dt = dt.predict(X_test)

print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))

... Decision Tree Accuracy: 0.9853658536585366
      precision    recall   f1-score   support
          0       0.97     1.00     0.99     102
          1       1.00     0.97     0.99     103

      accuracy                           0.99     205
   macro avg       0.99     0.99     0.99     205
weighted avg       0.99     0.99     0.99     205
```

```
[12] ✓ 0s ⏪ plt.figure(figsize=(4,4))
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Blues')
plt.title("Decision Tree Confusion Matrix")
plt.show()
```

A confusion matrix titled "Decision Tree Confusion Matrix". The matrix is 2x2, with classes 0 and 1. The diagonal elements are 102 (top-left) and 100 (bottom-right), both in dark blue. The off-diagonal elements are 0 (top-right) and 3 (bottom-left), both in light gray. A vertical color bar on the right indicates values from 0 to 100, with 100 being dark blue and 0 being light gray.

		0	1
0	102	0	
1	3	100	

```
[13] ✓ 0s ⏪ rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)

y_pred_rf = rf.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))

... Random Forest Accuracy: 0.9853658536585366
      precision    recall   f1-score   support
          0         0.97     1.00     0.99     102
          1         1.00     0.97     0.99     103

      accuracy                           0.99     205
   macro avg       0.99     0.99     0.99     205
weighted avg       0.99     0.99     0.99     205
```

```
[14] ✓ 0s ⏪ cm_rf = confusion_matrix(y_test, y_pred_rf)

plt.figure(figsize=(4,4))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Greens')
plt.title("Random Forest Confusion Matrix")
plt.show()

... Random Forest Confusion Matrix
      0         102        0
      1         3        100
      0         0         1
```

[15] ✓ 0s

```
▶ dt_params = {
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'criterion': ['gini', 'entropy']
}

dt_grid = GridSearchCV(
    DecisionTreeClassifier(random_state=42),
    dt_params,
    cv=5
)

dt_grid.fit(X_train, y_train)

print("Best Parameters for Decision Tree:", dt_grid.best_params_)

best_dt = dt_grid.best_estimator_
y_pred_dt_tuned = best_dt.predict(X_test)

print("Tuned Decision Tree Accuracy:",
      accuracy_score(y_test, y_pred_dt_tuned))

... Best Parameters for Decision Tree: {'criterion': 'gini', 'max_depth': None, 'min_samples_split': 2}
Tuned Decision Tree Accuracy: 0.9853658536585366
```

[16] ✓ 25s

```
▶ rf_params = {
    'n_estimators': [100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5]
}

rf_grid = GridSearchCV(
    RandomForestClassifier(random_state=42),
    rf_params,
    cv=5
)

rf_grid.fit(X_train, y_train)

print("Best Parameters for Random Forest:", rf_grid.best_params_)

best_rf = rf_grid.best_estimator_
y_pred_rf_tuned = best_rf.predict(X_test)

print("Tuned Random Forest Accuracy:",
      accuracy_score(y_test, y_pred_rf_tuned))

... Best Parameters for Random Forest: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 100}
Tuned Random Forest Accuracy: 0.9853658536585366
```

```
[17] ✓ 0s      accuracies = [
        accuracy_score(y_test, y_pred_dt_tuned),
        accuracy_score(y_test, y_pred_rf_tuned)
    ]

    plt.figure(figsize=(6,4))
    sns.barplot(x=models, y=accuracies)
    plt.title("Model Comparison (Accuracy)")
    plt.ylabel("Accuracy")
    plt.show()
```

